

Updated Last
5 Months Ago

Started In
March 2013

Gurobi.jl



[Gurobi.jl](#) is a wrapper for the [Gurobi Optimizer](#).

It has two components:

- a thin wrapper around the complete C API
- an interface to [MathOptInterface](#)

Affiliation

› This wrapper is maintained by the JuMP community with help from Gurobi.

Getting help

› If you are a commercial customer, please contact Gurobi directly through the [Gurobi Help Center](#).

Otherwise, you should ask a question on the [JuMP community forum](#), with the [gurobi](#) tag, or post in Gurobi's [Community Forum](#)

If you have a reproducible example of a bug, please [open a GitHub issue](#).

License

› [Gurobi.jl](#) is licensed under the [MIT License](#).

The underlying solver is a closed-source commercial product for which you must [obtain a license](#).

Free Gurobi licenses are available for [academics and students](#).

Installation

› To use Gurobi, you need a license, which you can obtain from [gurobi.com](#).

Once you have a license, follow Gurobi's instructions to [retrieve and set up a Gurobi license](#).

The instructions depend on the type of license that you have obtained.

As one exception, if you have used the default installation of Gurobi.jl and the instructions call for `grbgetkey`, do:

```
import Pkg
Pkg.add("Gurobi_jll")
import Gurobi_jll
# Replace the contents xxxxx with your actual key
key = "xxxxxxxx-xxxx-xxxx-xxxx-xxxxxxxxxxxxx"
run(`$(Gurobi_jll.grbgetkey()) $key`)
```

Default installation

› Install Gurobi as follows:

```
import Pkg
Pkg.add("Gurobi")
```

In addition to installing the Gurobi.jl package, this will also download and install the Gurobi binaries from [Gurobi_jll.jl](https://gurobi.com/Products/Software-Installation-Instructions/). You do not need to install Gurobi separately.

Manual installation

› To opt-out of using the `Gurobi_jll` binaries, set the `GUROBI_HOME` environment variable to point to your local installation and set the `GUROBI_JL_USE_GUROBI_JLL` environment variable to `"false"`, then run `Pkg.add` and `Pkg.build`:

```
# On Windows, this might be
ENV["GUROBI_HOME"] = "C:\\Program Files\\gurobi1100\\win64"
# ... or perhaps ...
ENV["GUROBI_HOME"] = "C:\\gurobi1100\\win64"
# On Mac, this might be
ENV["GUROBI_HOME"] = "/Library/gurobi1100/macos_universal2"

# Opt-out of using Gurobi_jll
ENV["GUROBI_JL_USE_GUROBI_JLL"] = "false"

import Pkg
Pkg.add("Gurobi")
Pkg.build("Gurobi")
```

To change the location of a manual install, change the value of `GUROBI_HOME`, re-run `Pkg.build("Gurobi")`, and then re-start Julia for the change to take effect.

Use with JuMP

› To use Gurobi with JuMP, use `Gurobi.Optimizer` :

```
using JuMP, Gurobi
model = Model{Gurobi.Optimizer}()
set_attribute(model, "TimeLimit", 100)
set_attribute(model, "Presolve", 0)
```

MathOptInterface API

› The Gurobi optimizer supports the following constraints and attributes.

List of supported objective functions:

- `MOI.ObjectiveFunction{MOI.ScalarAffineFunction{Float64}}`
- `MOI.ObjectiveFunction{MOI.ScalarQuadraticFunction{Float64}}`
- `MOI.ObjectiveFunction{MOI.VariableIndex}`
- `MOI.ObjectiveFunction{MOI.VectorAffineFunction{Float64}}`

List of supported variable types:

- `MOI.Reals`

List of supported constraint types:

- `MOI.ScalarAffineFunction{Float64}` in `MOI.EqualTo{Float64}`
- `MOI.ScalarAffineFunction{Float64}` in `MOI.GreaterThan{Float64}`
- `MOI.ScalarAffineFunction{Float64}` in `MOI.LessThan{Float64}`
- `MOI.ScalarQuadraticFunction{Float64}` in `MOI.EqualTo{Float64}`
- `MOI.ScalarQuadraticFunction{Float64}` in `MOI.GreaterThan{Float64}`
- `MOI.ScalarQuadraticFunction{Float64}` in `MOI.LessThan{Float64}`
- `MOI.VariableIndex` in `MOI.EqualTo{Float64}`
- `MOI.VariableIndex` in `MOI.GreaterThan{Float64}`
- `MOI.VariableIndex` in `MOI.Integer`
- `MOI.VariableIndex` in `MOI.Interval{Float64}`
- `MOI.VariableIndex` in `MOI.LessThan{Float64}`
- `MOI.VariableIndex` in `MOI.Semicontinuous{Float64}`
- `MOI.VariableIndex` in `MOI.Semiinteger{Float64}`
- `MOI.VariableIndex` in `MOI.ZeroOne`
- `MOI.VectorOfVariables` in `MOI.SOS1{Float64}`
- `MOI.VectorOfVariables` in `MOI.SOS2{Float64}`
- `MOI.VectorOfVariables` in `MOI.SecondOrderCone`

- [MOI.VectorAffineFunction](#) in [MOI.Indicator](#)

List of supported model attributes:

- [MOI.HeuristicCallback\(\)](#)
- [MOI.LazyConstraintCallback\(\)](#)
- [MOI.Name\(\)](#)
- [MOI.ObjectiveSense\(\)](#)
- [MOI.UserCutCallback\(\)](#)

Options

› See the [Gurobi Documentation](#) for a list and description of allowable parameters.

C API

› The C API can be accessed via `Gurobi.GRBxx` functions, where the names and arguments are identical to the C API.

See the [Gurobi documentation](#) for details.

As general rules when converting from Julia to C:

- When Gurobi requires the column index of a variable `x`, use `Gurobi.c_column(model, x)`
- When Gurobi requires a `Ptr{T}` that holds one element, like `double *`, use a `Ref{T}()`.
- When Gurobi requires a `Ptr{T}` that holds multiple elements, use a `Vector{T}`.
- When Gurobi requires a `double`, use `Cdouble`
- When Gurobi requires an `int`, use `Cint`
- When Gurobi requires a `NULL`, use `C_NULL`

For example:

```
julia> import MathOptInterface as MOI

julia> using Gurobi

julia> model = Gurobi.Optimizer();

julia> x = MOI.add_variable(model)
MOI.VariableIndex(1)

julia> x_col = Gurobi.c_column(model, x)
0

julia> GRBupdatemodel(model)
```

0

```
julia> pValue = Ref{Cdouble}(NaN)
Base.RefValue{Float64}(NaN)
```

```
julia> GRBgetdblattrelement(model, "LB", x_col, pValue)
0
```

```
julia> pValue[]
-1.0e100
```

```
julia> GRBsetdblattrelement(model, "LB", x_col, 1.5)
0
```

```
julia> GRBupdatemodel(model)
0
```

```
julia> GRBgetdblattrelement(model, "LB", x_col, pValue)
0
```

```
julia> pValue[]
1.5
```

The C API from JuMP

- › You can call the C API from JuMP if you use `direct_model`. This is most useful for adding `GRBaddgenXXX` constraints. Here are some examples:

```
using JuMP, Gurobi
column(x::VariableRef) = Gurobi.c_column(backend(owner_model(x)), index(x))
model = direct_model(Gurobi.Optimizer())
@variable(model, x)
@variable(model, y)
p = [3.0, 0.0, 0.0, 7.0, 3.0]
GRBaddgenconstrPoly(backend(model), C_NULL, column(x), column(y), 5, p, "")
optimize!(model)
```

```
using JuMP, Gurobi
column(x::VariableRef) = Gurobi.c_column(backend(owner_model(x)), index(x))
model = direct_model(Gurobi.Optimizer())
@variable(model, x[i in 1:2])
@variable(model, y[1:2])
GRBaddgenconstrPow(backend(model), "x1^0.7", column(x[1]), column(y[1]), 0.7, "")
GRBaddgenconstrPow(backend(model), "x2^3", column(x[2]), column(y[2]), 3.0, "")
@objective(model, Min, y[1] + y[2])
optimize!(model)
```

Reusing the same Gurobi environment for multiple solves

- When using this package via other packages such as [JuMP.jl](#), the default behavior is to obtain a new Gurobi license token every time a model is created. If you are using Gurobi in a setting where the number of concurrent Gurobi uses is limited (for example, ["Single-Use"](#) or ["Floating-Use" licenses](#)), you might instead prefer to obtain a single license token that is shared by all models that your program solves.

You can do this by passing a `Gurobi.Env()` object as the first parameter to `Gurobi.Optimizer`. For example:

```
using JuMP, Gurobi
const GRB_ENV = Gurobi.Env()

model_1 = Model(() -> Gurobi.Optimizer(GRB_ENV))

# The solvers can have different options too
model_2 = direct_model(Gurobi.Optimizer(GRB_ENV))
set_attribute(model_2, "OutputFlag", 0)
```

If you create a module with a `Gurobi.Env` as a module-level constant, use an `__init__` function to ensure that a new environment is created each time the module is loaded:

```
module MyModule

import Gurobi

const GRB_ENV_REF = Ref{Gurobi.Env}()

function __init__()
    global GRB_ENV_REF
    GRB_ENV_REF[] = Gurobi.Env()
    return
end

# Note the need for GRB_ENV_REF[] not GRB_ENV_REF
create_optimizer() = Gurobi.Optimizer(GRB_ENV_REF[])

end
```

Accessing Gurobi-specific attributes

- Get and set Gurobi-specific variable, constraint, and model attributes as follows:

```
using JuMP, Gurobi
model = direct_model(Gurobi.Optimizer())
```

```

@variable(model, x >= 0)
@constraint(model, c, 2x >= 1)
@objective(model, Min, x)
grb = backend(model)
MOI.set(grb, Gurobi.ConstraintAttribute("Lazy"), index(c), 2)
optimize!(model)
MOI.get(grb, Gurobi.VariableAttribute("LB"), index(x)) # Returns 0.0
MOI.get(grb, Gurobi.ModelAttribute("NumConstrs")) # Returns 1

```

A complete list of supported Gurobi attributes can be found in [their online documentation](#).

Callbacks

Here is an example using Gurobi's solver-specific callbacks.

```

using JuMP, Gurobi, Test

model = direct_model(Gurobi.Optimizer())
@variable(model, 0 <= x <= 2.5, Int)
@variable(model, 0 <= y <= 2.5, Int)
@objective(model, Max, y)
cb_calls = Cint[]
function my_callback_function(cb_data, cb_where::Cint)
    # You can reference variables outside the function as normal
    push!(cb_calls, cb_where)
    # You can select where the callback is run
    if cb_where != GRB_CB_MIPSOL && cb_where != GRB_CB_MIPNODE
        return
    end
    # You can query a callback attribute using GRBcbget
    if cb_where == GRB_CB_MIPNODE
        resultP = Ref{Cint}()
        GRBcbget(cb_data, cb_where, GRB_CB_MIPNODE_STATUS, resultP)
        if resultP[] != GRB_OPTIMAL
            return # Solution is something other than optimal.
        end
    end
    # Before querying `callback_value`, you must call:
    Gurobi.load_callback_variable_primal(cb_data, cb_where)
    x_val = callback_value(cb_data, x)
    y_val = callback_value(cb_data, y)
    # You can submit solver-independent MathOptInterface attributes such as
    # lazy constraints, user-cuts, and heuristic solutions.
    if y_val - x_val > 1 + 1e-6
        con = @build_constraint(y - x <= 1)
        MOI.submit(model, MOI.LazyConstraint(cb_data), con)
    elseif y_val + x_val > 3 + 1e-6
        con = @build_constraint(y + x <= 3)
        MOI.submit(model, MOI.LazyConstraint(cb_data), con)
    end
end

```

```

    if rand() < 0.1
        # You can terminate the callback as follows:
        GRBterminate(backend(model))
    end
    return
end
# You _must_ set this parameter if using lazy constraints.
MOI.set(model, MOI.RawOptimizerAttribute("LazyConstraints"), 1)
MOI.set(model, Gurobi.CallbackFunction(), my_callback_function)
optimize!(model)
@test termination_status(model) == MOI.OPTIMAL
@test primal_status(model) == MOI.FEASIBLE_POINT
@test value(x) == 1
@test value(y) == 2

```

See the [Gurobi documentation](#) for other information that can be queried with `GRBcbget`.

Common Performance Pitfall with JuMP

- › Gurobi's API works differently than most solvers. Any changes to the model are not applied immediately, but instead go sit in a internal buffer (making any modifications appear to be instantaneous) waiting for a call to `GRBupdatemodel` (where the work is done).

This leads to a common performance pitfall that has the following message as its main symptom:

```
Warning: excessive time spent in model updates. Consider calling update less frequently.
```

This often means the JuMP program was structured in such a way that `Gurobi.jl` ends up calling `GRBupdatemodel` in each iteration of a loop.

Usually, it is possible (and easy) to restructure the JuMP program in a way it stays solver-agnostic and has a close-to-ideal performance with Gurobi.

To guide such restructuring it is good to keep in mind the following bits of information:

1. `GRBupdatemodel` is only called if changes were done since last `GRBupdatemodel` (that is, if the internal buffer is not empty).
2. `GRBupdatemodel` is called when `JuMP.optimize!` is called, but this often is not the source of the problem.
3. `GRBupdatemodel` *may* be called when *any* model attribute is queried, *even if* that specific attribute was not changed. This often the source of the problem.

The worst-case scenario is, therefore, a loop of modify-query-modify-query, even if what is being modified and what is being queried are two completely distinct things.

As an example, instead of:


```

model = Model(Gurobi.Optimizer)
@variable(model, x[1:100] >= 0)
for i in 1:100
    set_upper_bound(x[i], i)
    # `GRBupdatemodel` called on each iteration of this loop.
    println(lower_bound(x[i]))
end

```

do

```

model = Model(Gurobi.Optimizer)
@variable(model, x[1:100] >= 0)
# All modifications are done before any queries.
for i in 1:100
    set_upper_bound(x[i], i)
end
for i in 1:100
    # Only the first `lower_bound` query may trigger an `GRBupdatemodel`.
    println(lower_bound(x[i]))
end

```

Common errors

› Using Gurobi v9.0 and you got an error like `q not PSD ?`

› You need to set the NonConvex parameter:

```

model = Model(Gurobi.Optimizer)
set_optimizer_attribute(model, "NonConvex", 2)

```

Gurobi Error 1009: Version number is XX.X, license is for version XX.X

› Make sure that your license is correct for your Gurobi version. See the [Gurobi documentation](#) for details.

Once you are sure that the license and Gurobi versions match, re-install Gurobi.jl by running:

```

import Pkg
Pkg.build("Gurobi")

```

Required Packages

- [BenchmarkTools](#)